

COM3110/4115/6115:

Text Processing

Text Compression

Rob Gaizauskas

Department of Computer Science
University of Sheffield

Overview

- Models
 - ◊ Static
 - ◊ Semi-static
 - ◊ Adaptive
- Coding
 - ◊ Huffman Coding
 - ◊ Arithmetic Coding
- Further topics:
 - ◊ Symbolwise Models
 - ◊ Dictionary Methods
 - ◊ Synchronisation
 - ◊ Performance Issues

Introduction

- Have seen a dramatic increase of
 - ◇ low cost disk storage
 - ◇ transmission bandwidth
 - ◇ processor speed
- But also a massive increase in data volume:
 - ◇ text, sound and images
 - ◇ so, techniques to *compress* data remain significant
- We shall concentrate on techniques for *text compression*
- Text compression distinct from some other forms of data compression:
 - ◇ the text must be *exactly reconstructable*
 - ◇ not so critical for digitised analogue signals, such as image/sound
 - ◇ text compressions requires so-called *lossless coding*

Introduction (ctd)

- Distinguish *lossless* vs. *lossy* compression methods
- **Lossless** Compression:
 - ◇ class of algorithms allowing original data to be perfectly reconstructed from compressed data
- **Lossy** Compression:
 - ◇ achieve data reduction by *discarding* (i.e. losing) information
 - ◇ suitable for certain media types, esp.: image / video / audio data
 - widely used in data *streaming* contexts
 - e.g. achieve data reduction of an image by computing a version with lower pixel density
 - ◇ Text data requires *lossless* compression
 - “text from which $N\%$ of info discarded” doesn’t make sense
 - expect decompression to return text identical to original in form/content

(See Wikipedia pages for Lossy and Lossless Compression)

Introduction (ctd)

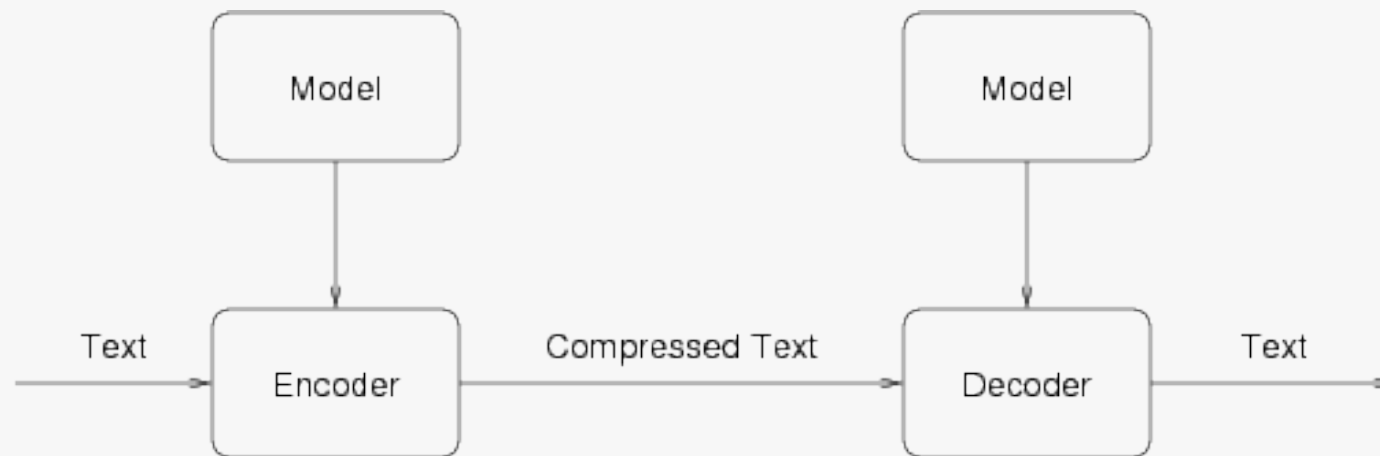
- Text compression techniques may be classified in several ways.
 - ◇ *symbolwise* vs. *dictionary* methods
 - ◇ *static* vs. *adaptive* methods
- **Dictionary methods** work by replacing word/text fragments with an index to an entry in a dictionary
- **Symbolwise methods** work by estimating the probabilities of symbols (characters/words) and coding one symbol at a time using shorter codewords for the more likely symbols
 - ◇ rely on a **modeling** step and a **coding** step
 - ◇ *modeling*: estimation of probabilities for the symbols in the text
 - the better the probability estimates, the higher the compression that can be achieved
 - ◇ *coding*: conversion of probabilities from model into a bitstream

Introduction (ctd)

- **Static** vs. **adaptive** methods:
 - ◇ *Static*: use a fixed model or fixed dictionary derived in advance of any text to be compressed
 - ◇ *Semi-static*: use current text to build a model or dictionary during one pass, then apply it in second pass
 - ◇ *Adaptive*: build model or dictionary adaptively during one pass

Coding and Decoding

- Function of model is to predict symbols
 - ◇ model amounts to a probability distribution for all possible symbols, i.e. the “alphabet”
- The encoder uses the model to encode (compress) the text
- The decoder must use the same model to decode it



Information Content

- The number of *bits* in which a symbol s should be coded is its *information content*, denoted $I(s)$
- Information content related to *predicted probability* $P[s]$, as follows:

$$I(s) = -\log_2 P[s]$$

e.g. for a fair coin toss in which the outcome is “heads”, the best an encoder can do it use $-\log_2(\frac{1}{2}) = 1$ bit.

- The *entropy* H of the probability distribution — the average information per symbol over the whole alphabet — is given by:

$$H = \sum_s P[s] \cdot I(s) = \sum_s -P[s] \cdot \log_2 P[s] \quad \text{bits/character}$$

◇ akin to ‘average’ of code lengths *weighted* by probability

- The entropy H places a *lower bound* on compression
 - ◇ Shannon’s Source Coding Theorem

Information Content (ctd)

Examples:

- With a 'fair' coin, $P(\text{head}) = P(\text{tail}) = 0.5$
 - ◇ hence: $I(\text{head}) = I(\text{tail}) = -\log_2(0.5) = 1$
 - ◇ hence: $H = 0.5 \cdot 1 + 0.5 \cdot 1 = 1$
- With a 'biased' coin, such that $P(\text{head}) = 0.99$, $P(\text{tail}) = 0.01$
 - ◇ $I(\text{head}) = -\log_2(0.99) = 0.0145$, $I(\text{tail}) = -\log_2(0.01) = 6.644$
 - ◇ $H = 0.99 \cdot 0.0145 + 0.01 \cdot 6.644 = 0.0808$
- For a 'fair' 8-sided dice, $P(s) = \frac{1}{8} = 0.125$ for each side s
 - ◇ $I(s) = -\log_2(\frac{1}{8}) = 3$, and hence $H = 3$ also
- For a 'biased' 8-sided dice, with $P(s_0) = 0.9$ and other $P(s_i) = 0.0143$
 - ◇ $I(s_0) = -\log_2(0.9) = 0.152$, $I(s_i) = -\log_2(0.0143) = 6.13$
 - ◇ $H = 0.7498$

Models and Context

- Probability of encountering a given symbol at a particular place in a text is influenced by *preceding symbols*
e.g. probability of u following q much higher than u occurring on average
- Models that take immediately preceding symbols into account are called *finite-context models*
 - ◇ best text compression results consider contexts of 3-5 characters
- Models that *ignore* preceding content are called *zero order models*

Static Zero Order Character-based Model for Moby Dick

Ch	Count	Pr	Ch	Count	Pr	Ch	Count	Pr	Ch	Count	Pr
SP	198111	0.1623	p	16207	0.0132	C	1148	0.0009	K	178	0.0001
e	115863	0.0949	b	15453	...	P	1049	...	V	171	...
t	85544	0.0701	v	8428		x	1008		l	140	
a	75267	...	k	7882		?	1004		0	131	
o	68341		.	7558		O	990		2	60	
n	64434		-	5984		L	901		8	58	
s	62023		;	4174		j	830		5	54	
i	61893		I	3544		R	824		7	53	
h	61435		"	3071		F	804		3	47	
r	751314		'	2922		M	754		*	45	
l	741896		A	2651		D	752		4	39	
d	37469		T	2458		G	641		Z	38	
u	26458		S	2209		z	594		6	37	
NL	22933		!	1767		Y	331		9	35	
m	22526		H	1465		Q	322		-	26	
c	21361		B	1427		J	253		X	23	
w	20917		W	1306		U	240		&	2	
g	20181		E	1240		(215		\$	2	
f	20031		q	1234)	215		[2	
,	19230		N	1186		:	196]	2	1.6e-06
y	16543										

- ◇ 81 characters in alphabet
- ◇ total character occurrences = 1,220,150
- ◇ entropy using this model: 4.4953 bits/char

Static vs. Adaptive Modelling

Probabilities may be estimated in various ways

- **Static modelling** derives, and then uses, a single model for all texts
 - ◇ will perform poorly on texts different from those used in constructing the model, e.g. texts with tables of numbers
- **Semi-static modelling** derives model for the file in a 1st pass
 - ◇ model derived will be better suited to the text than a static one
 - ◇ but, is inefficient because:
 - must make two passes over text
 - must also transmit model
- **Adaptive modelling** derives model during encoding

Adaptive Models

- Adaptive models begin with a *base* probability distribution
 - ◇ refine the model as more symbols are encountered, during encoding
 - ◇ so, the text being encoded itself *re-defines the model*
- Decoder starts with same base probability distribution
 - ◇ it is decoding the *same symbol sequence*
 - ◇ so, it can refine the model in the same way
- *Issues:*
 - ◇ Care must be taken to ensure no character is ever predicted with zero probability simply because it has not been seen yet
 - ◇ Principal disadvantage: not suitable for random access to files
 - a text can only be decoded from the beginning
 - poses difficulties for retrieval applications

Adaptive Models (ctd)

- *EXAMPLE:* encoding Moby Dick
 - ◇ might start assuming a uniform probability of $1/81$ per char
 - implies same minimal possible code length for all symbols at $-\log_2(1/81) = 6.34$ bits
 - ◇ after encoding (say) 100,000 chars, and observing char distributions, we have a much more accurate 0-order model
 - ◇ suppose now about to encode e of whale:
 - by now have observed that $\sim 10\%$ of chars are e
 - thus, the e can be encoded in $\sim -\log_2 0.1 = 3.32$ bits
 - decoder will refine its model in parallel manner for use in decoding


Huffman Coding

- **Coding** is the task of determining the output representation of a symbol, given the probability distribution supplied by the model
- Coder should output:
 - ◇ short codes for high probability symbols
 - ◇ long codes for low probability symbols
- Speed of coder may also be significant
 - ◇ computing optimal codes can be slow
 - ◇ hence, there is a trade-off between speed of compression and compression rate
- Huffman coding dates from 1952
 - ◇ was the dominant model till the 1970's
 - ◇ with refinements, it still has applications today, e.g. in text retrieval

Huffman Coding (ctd)

- Technique uses a **code tree** for encoding and decoding
 - ◇ each branch is labelled with a 0 or 1
 - ◇ each leaf node is a symbol in the alphabet
 - ◇ code is **prefix-free** – no codeword is the prefix of another
- Consider a seven symbol alphabet (example from Witten *et al.*)

Symbol	Prob	Codeword
a	0.05	0000
b	0.05	0001
c	0.1	001
d	0.2	01
e	0.3	10
f	0.2	110
g	0.1	111



e.g. eefggfed is coded as 10101101111111101001

Huffman Coding (ctd)

- Given a set of codewords produced by the algorithm, can compute the expected *average* code length as follows:

- ◇ multiply each symbol code length by associated probability
- ◇ sum results across all symbols

Symbol	Prob	Code	len	$p \times len$
a	0.05	0000	4	0.2
b	0.05	0001	4	0.2
c	0.1	001	3	0.3
d	0.2	01	2	0.4
e	0.3	10	2	0.6
f	0.2	110	3	0.6
g	0.1	111	3	0.3
				2.6

- Compare this to minimal *fixed-length* code length for same symbol set
e.g. for above symbol set, could use a 3-bit fixed length code
 - ◇ 3-bit fixed length code allows for $2^3 = 8$ symbols


Huffman — the coding algorithm

- The code tree is constructed bottom up from the probabilistic model according to the following algorithm:
 - (a) Probabilities are associated with leaf nodes
 - (b) Identify the two nodes with smallest probabilities
 - join them under a parent node, whose probability is their sum
 - (c) Repeat step (b) until only one node remains
 - (d) 0's and 1's are then assigned to each binary split
- *EXAMPLE:*

- ◇ start with leaf nodes:

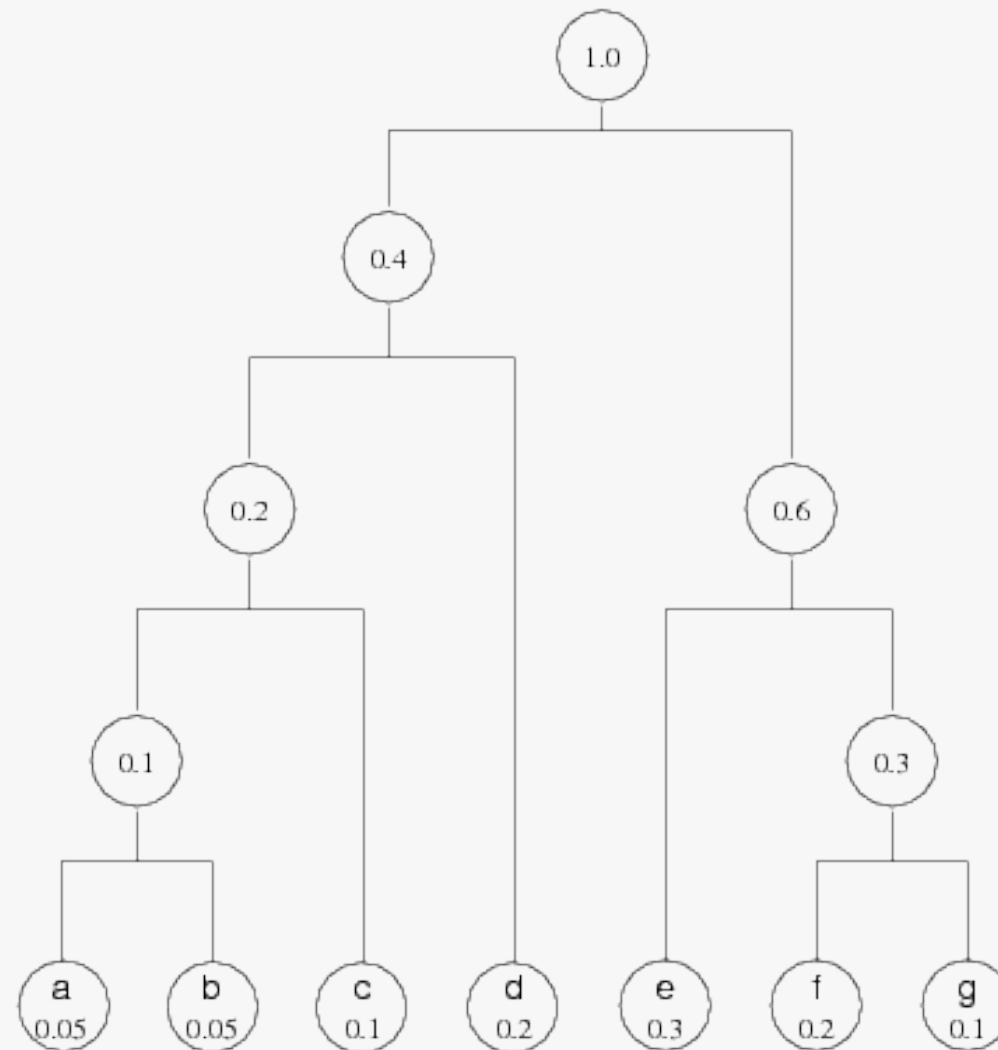


- ◇ join two nodes with smallest probabilities:



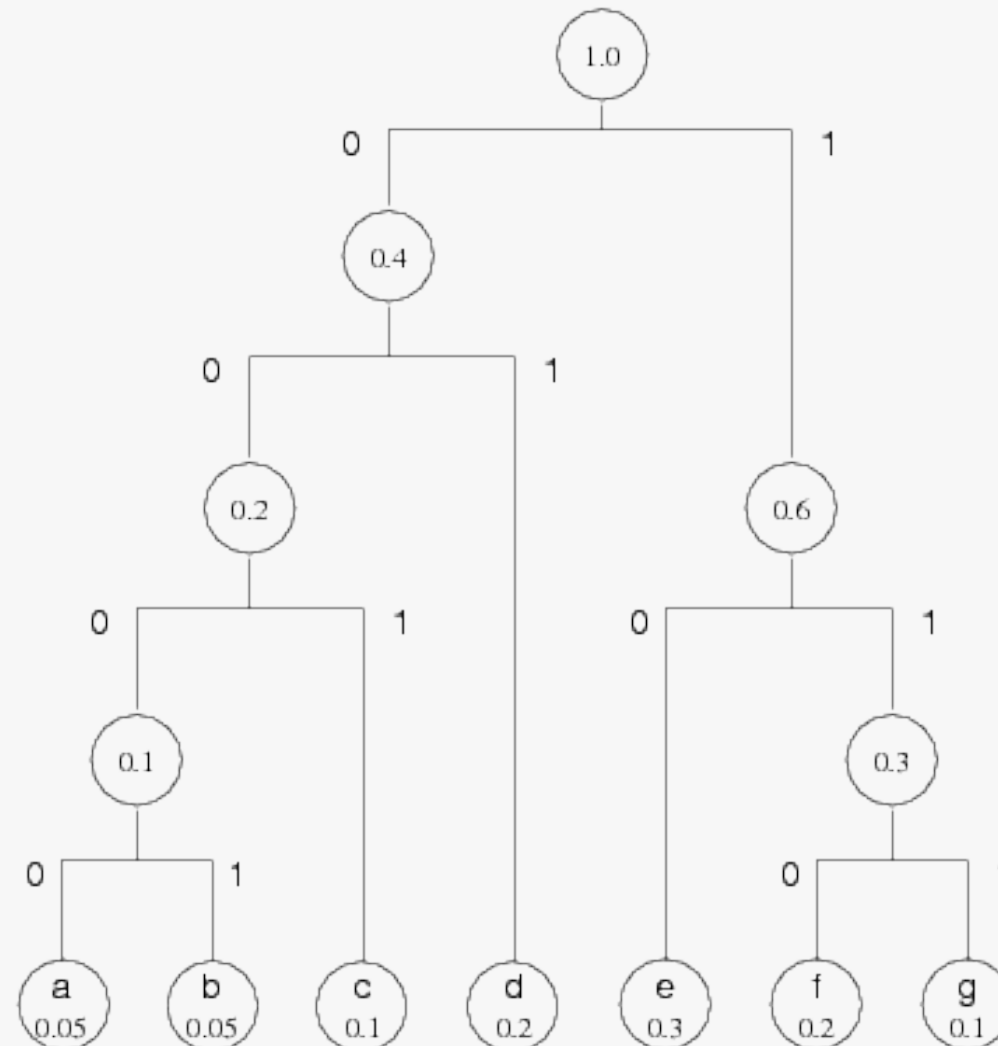
Huffman — the coding algorithm (ctd)

- ◇ repeat until only one node remains ...



Huffman — the coding algorithm (ctd)

- ◇ assign 0/1 to branches of each binary fork
 - doesn't matter which is which



Huffman Coding — further details

- Huffman coding is fast for both encoding and decoding
 - ◇ provided probability distribution is static
- Variants of Huffman coding developed for adaptive models
 - ◇ *but* these are complex
 - ◇ generally better to use arithmetic coding for adaptive models
- Huffman effective when used with a word-based model, rather than a character-based model
 - ◇ gives good compression
 - ◇ fast
 - ◇ supports random access to compressed files
(given some additional requirements on how method used)

Canonical Huffman codes

- Huffman gives effective compression where there are *many symbols*, with a highly *skewed* distribution
 - ◇ arises, e.g. for *word-based* models
 - ◇ i.e. many symbols (words), some v.frequent, some rare
- *Tree-based* storage of *v.large models* is *costly & inefficient*
 - ◇ *tree nodes* store pointers to *child nodes* — *costly for memory*
 - ◇ *traversing* tree involves much *jumping* between locations — *inefficient*
- These issues addressed by use of *canonical Huffman codes*
 - ◇ special Huffman code, where codes generated in a *standardised format*
 - ◇ specifically, all codes for given code-length *assigned values sequentially*
 - ◇ this feature allows for both
 - efficient storage and/or transmission of codebook
 - more efficient decoding algorithm

Canonical Huffman codes (ctd)

- To create a canonical code:
 - ◇ first determine length of code for each symbol
 - can do this by applying standard Huffman coding algorithm
 - ◇ group symbols having same code-length into blocks, and order
 - e.g. could sort alphabetically, or in some other way
 - ◇ assign codes by '*counting up*' – addressing blocks in code-length order

- To illustrate, consider our 7-letter example again:

Symbol	Prob	Code	<i>len</i>
a	0.05	0000	4
b	0.05	0001	4
c	0.1	001	3
d	0.2	01	2
e	0.3	10	2
f	0.2	110	3
g	0.1	111	3

- ◇ ignore original codes, and group symbols by their code-length, i.e.:

[2] d e [3] c f g [4] a b

Canonical Huffman codes (ctd)

- *Example – continued*: creating a canonical code ...

- ◊ First, group symbols by their code lengths:

[2] d e [3] c f g [4] a b

- ◊ Assign codes:

- assign first symbol a 'zero' code of required length
- for successive symbols, simply count up 1
- if code length goes up, then (i) count up & (ii) add 0s to get new len

[2]	d	e	[3]	c	f	g	[4]	a	b
	00	01		100	101	110		1110	1111

- To store / transmit this code, is sufficient to specify:

- ◊ sequence of symbols: d e c f g a b
- ◊ number of symbols at each code-length: (0, 2, 3, 2)
 - i.e. *len 1: 0 items; len 2: 2 items; len 3: 3 items; len 4: 2 items*
- ◊ this is sufficient info to reconstruct entire code

Canonical Huffman codes (ctd)

- This approach also supports an efficient decoding algorithm
 - ◇ does not require a code tree
- Method:
 - ◇ store blocks of symbols in sequential order
 - ◇ compute code for only the *first* symbol in each block
 - ◇ by comparing prefixes of input to these first-symbol codes ($<$, $>$), can determine which block next code belongs to, and hence its length
 - ◇ binary difference between next code and first-symbol code gives position of symbol in block sequence
- *Example:*
 - ◇ Example on last slide has first-symbol codes: 00, 100, 1110
 - ◇ Assume input: 11001111.....
 - ◇ Find: prefix 11 $>$ 00; prefix 110 $>$ 100; prefix 1100 $<$ 1110
 - so next codeword is 110, of the length 3 block
 - ◇ Difference of 100 and 110 shows symbol is 3rd item in block sequence

Reading

- Main:
 - ◇ Baeza-Yates and Ribeiro-Neto, Ch 7.4-7.5
- Other:
 - ◇ I. H. Witten, A. Moffat, T. C. Bell. Managing Gigabytes: Compressing and Indexing Documents and Images, 2nd ed. Morgan Kaufmann. 1999.
 - ◇ Nam Phamdo. Theory of Data Compression.
www.data-compression.com/theory.html